

## 05 SCALING STRATEGY AND ROADMAP

---

### Phase 1 - Production Hardening

---

Replace SQLite with PostgreSQL and formalize schema migrations. Add environment-specific config profiles and secure secret management.

Introduce robust authentication/authorization boundaries and role-scoped API guards where currently permissive for demo speed.

Add observability foundation: request IDs, structured logs, endpoint latency metrics, model call stats, and error alerting.

### Phase 2 - Throughput and Reliability

---

Move expensive tasks (resume parsing, report regeneration) to async workers with queue-based retries.

Add object storage for large artifacts (resume files, optional evidence snapshots) and keep DB for metadata/indexing.

Implement cache and invalidation rules for expensive aggregate reads to protect recruiter dashboard responsiveness at scale.

### Phase 3 - Enterprise Controls

---

Add SSO/SAML and granular RBAC.

Implement policy editor for baseline rules per role/department.

Support multi-tenant boundaries with explicit tenant IDs and isolation guardrails.

Add compliance exports and immutable audit trail generation for decision records.

### Phase 4 - Model Governance

---

Track model output quality against hiring outcomes and false positive/negative rates.

Enable controlled prompt/version rollouts and A/B validation for recommendation quality.

Add confidence calibration using historical recruiter decisions and post-hire performance proxies.

### What Could Be Achieved with More Time

---

A policy-governed hiring intelligence platform where recommendation quality is measurable, explainable, and operationally auditable across teams and roles.

A recruiter workflow that moves from manual subjective reviews to signal-based decision support with clear override controls.