

System Architecture Questions

Q: Why this stack? A: FastAPI and TypeScript provide typed interfaces on both backend and frontend, minimizing contract drift and enabling rapid iteration with strong maintainability.

Q: Why not microservices? A: For PoC speed and coherence, modular monolith was chosen. Service boundaries already exist in code and can be split later with lower risk.

Q: How do frontend and backend stay in sync? A: Typed API wrappers in frontend and Pydantic models in backend keep request/response contract explicit.

AI and Explainability Questions

Q: Is your AI black-box? A: No. Output is constrained to strict sections and evidence tags, and fallback logic prevents uncontrolled behavior.

Q: How do you avoid unnecessary AI cost? A: Rationale regeneration is controlled by input-state checks and reuse logic where applicable.

Q: How do you handle model outages? A: Rule-based fallback generates recommendation and rationale structure without external dependency.

Integrity and Security Questions

Q: Can candidates cheat camera checks? A: Camera is supplementary. Deterministic telemetry like tab-switch and paste events are weighted and integrated into recommendation.

Q: Why trust your integrity score? A: It is explicitly computed from persisted severity events and visible in recruiter dashboard; it is not hidden model output.

Q: Are false positives possible? A: Yes, especially camera anomalies. The system mitigates this by multi-signal evaluation and review-capable outcomes.

Scalability and Cost Questions

Q: What breaks first at scale? A: SQLite and synchronous heavy operations. Roadmap already includes PostgreSQL and async worker separation.

Q: How expensive is model usage? A: Cost is bounded by controlled invocation points and reuse/fallback strategies.

Q: How would you make this enterprise-ready? A: Add SSO, RBAC, tenant isolation, observability, and policy governance layer.

Product and Evaluation Questions

Q: Why better than existing ATS? A: We provide traceable recommendation evidence and explicit role baseline checks, not only ranking or keyword scoring.

Q: What if candidate has good resume but poor test? A: Consistency risk is surfaced and recommendation can degrade accordingly.

Q: What if candidate has good test but poor integrity? A: Integrity can override outcome based on severity/volume policy.